⚠️ **ARCHIVED – Static Release**

The "**llm2-profile-search-audit-pm-report**" is a final static reference on 2026-04-23. This document is no longer actively maintained or updated.

LEGAL DISCLAIMER, WAIVER & DISCLOSURE

© ZZZ_EPOCHE 2026 | MIT License

This is an independent open-source research tool and audit artifact. The author has no affiliation with any large language model (LLM) provider, AI laboratory, or commercial entity.

This report is released publicly under the MIT License solely as an open-source transparency and AI safety research artifact. It is intended to contribute to the broader community discussion on LLM reliability, groundedness, and human oversight. Any republication, modification, or derivative work must retain this full disclaimer and legal waiver.

EU/EEA Warning

This report and the underlying methodology are **not intended** for use in the European Union or European Economic Area (EEA) as part of any high-risk AI system as defined under the EU AI Act. It does not constitute a conformity assessment, risk management system, technical documentation, or any other regulated compliance deliverable. Users in the EU/EEA must ensure full compliance with the EU AI Act and all applicable laws and regulations before any use or deployment. This report itself is **not** technical documentation, risk assessment, or post-market monitoring under the EU AI Act. Users assume all risks associated with any use, citation, or deployment decisions based on this report.

Intended Use

This post-mortem report is provided strictly for **academic research**, **defensive safety engineering**, **responsible AI governance experimentation**, **self-audit practices**, and other **non-commercial defensive purposes** only.

It serves as a transparency and learning artifact to document, analyze, and remediate errors in LLM behavior and should not be interpreted as formal legal advice, security certification, or production-grade audit.

"AS IS" Provision

This report is provided **"AS IS"** and **"AS AVAILABLE"**, without any warranties of any kind, express or implied, including but not limited to the warranties of merchantability, fitness for a particular purpose, accuracy, completeness, reliability, or non-infringement.

No Guarantees

- This post-mortem enumerates identified mistakes and proposed remediations but **does not guarantee** completeness, correctness, or that all errors have been captured.
- Error classification, severity ratings, root cause analysis, and remediation suggestions are based on limited review and may contain omissions, inaccuracies, or subjective judgments.
- All performance claims, capability assessments, and compliance matrices are derived from internal analysis only and are **not** independently verified third-party audit results.
- The report does **not** guarantee prevention of future mistakes or full compliance with any external standard.

Liability Waiver

To the fullest extent permitted by applicable law, by using, reading, citing, modifying, or relying on this Report in any way, you expressly waive, release, and discharge the author(s) from any and all claims, liabilities, damages, or losses (including direct, indirect, consequential, special, punitive, or exemplary damages) arising from its use or reliance, even if the author(s) have been advised of the possibility of such damages.

All use and reliance is strictly at **your own risk and sole responsibility.**

User Responsibility

You are solely responsible for:

- Thoroughly reviewing, validating, and independently verifying all findings, error tallies, root causes, and remediation steps
- Ensuring full compliance with all applicable local, national, and international laws, regulations (including the EU AI Act where relevant), and ethical standards
- All decisions and actions taken based on this report
- Any consequences resulting from the use of this document or its recommendations

Commercial use or integration into commercial products/services is **strictly prohibited** without prior written permission from the author.

Static Release

This is a static release. No ongoing maintenance, updates, corrections, or technical support is currently planned or provided.

No-Endorsement / No-Defamation

This report contains an independent analysis of specific LLM responses and does not constitute defamation, disparagement, or any intentional harm toward any AI provider, model, or company. All findings are presented in good faith for transparency, safety research, and educational purposes only. The author makes no claims about the overall quality, safety, or performance of any commercial LLM beyond the documented thread.

Evaluations and Matrices Bias Waiver and Disclaimer
All evaluation scores, graphs, estimations, matrices, comparisons, and qualitative assessments in this report are the result of **internal**, **subjective analysis** conducted solely for transparency and educational purposes only. They reflect the author's personal weighting, design goals, limited interactions with the specific thread, and internal simulations. They inherently carry **bias** toward the project's emphasis on human sovereignty, operator control, epistemic rigor, and upfront honesty.
These metrics **do not** represent independent third-party benchmarks, production-grade evaluations, or statistically validated data. No external validation or certification has been performed.
The author explicitly disclaims any and all liability arising from reliance on these evaluations, scores, graphs, or matrices. Users bear sole responsibility for any decisions, interpretations, or actions taken based on this report and must conduct their own independent verification and due diligence.
No private, confidential, or personally identifiable information (PII) was disclosed, processed, or referenced in the conversation thread or in this sanitized post-mortem report. All user interactions were handled with full respect for privacy, and the report has been fully anonymized and sanitized.
By accessing, reading, using, or relying on this Report, in any way, you acknowledge that you have read, understood, and fully agree to all terms of this Legal Disclaimer, Waiver, and Disclosure.

Name: llm2-profile-search-audit-pm-report

Version: v1.0 – Final Post-Mortem Report (Static release)

Report ID: llm2-profile-search-audit-pm-report-20260423

Code Name: llm2-p-s-audit-pm-report-20260423

Format: PDF

Date of release: 2026-04-23

Author: ZZZ_EPOCHE

Assisted by: Frontier LLM under outer-governance invariants

AI Assistance Disclosure: Portions of this report were drafted or refined with assistance from one or more AI language models under outer-governance invariants. The author remains fully responsible for the content, accuracy, and final wording of all sections.

Detection Provenance Initiated by: user

Executive Summary

This report reviews LLM2's responses to a user request for GitHub profile analysis. It identifies 15 distinct issues (7 High, 6 Medium, 2 Low severity), mainly related to capability representation and grounding/accuracy gaps. The primary issues arose from offering structured analysis without first clearly stating tool and access limitations.

Key findings include gaps in factual grounding, inaccurate descriptions of capabilities, and shortcomings in transparency. The report outlines contributing factors, observed positive behaviors during correction, and suggested improvements. Its goal is to support clearer communication of capabilities and reduce similar issues in future interactions.

- Main observed patterns: tendency toward decisive statements about capabilities, verification gaps, and unlabeled assumptions.
- Key takeaway: Human oversight remains important for maintaining accuracy; LLMs can prioritize conversational flow over strict verification when tool constraints apply.
- Impact: Accumulation of inaccuracies that reduced user confidence and led to session termination.

Scope of This Report

This document serves as a structured review of a specific conversation thread with LLM2. It consolidates identified issues, examines contributing factors, and suggests practical improvements. It is intended for AI reliability learning, self-audit practice, and general reference in responsible AI development.

Limitations of This Report

- This is not a formal security audit, compliance certification, or regulatory assessment.
- It does not guarantee completeness or error-free analysis.
- It does not provide legal advice or formal risk documentation.

- It does not constitute an official evaluation of any AI provider or model beyond the reviewed thread.

Potential Future Improvements

This lightweight report format could be refined over time through better tooling and automation. Its current strengths include transparency, speed, and low overhead. Limitations include reliance on manual effort and limited scalability for enterprise or high-risk uses. The focus remains on practical learning and openness rather than production-grade robustness.

Contents of This Report

- Disclaimers and limitations
- Detailed issue enumeration with severity and recommendations
- Severity distribution and summary tables
- Remediation checklist and updated protocols
- Suggestions for clearer tool use and citation practices

EVALUATION MATRICES *Evaluated on a 10-point scale.*

Table 1: Report & Self-Audit Methodology Evaluation.

Criterion	Score	Notes
Cleanliness & Readability	8.5	Clear structure, consistent formatting, use of sections and tables
Modularity	8.0	Well-organized into disclaimer, errors, root causes, matrices, and closing
Forensic Logging	7.5	Strong provenance tracking and error enumeration; could add more timestamps
Token Efficiency	9.0	Concise yet thorough; minimal fluff
Safety Layering	8.5	Disclaimers, privacy protections, and human oversight focus

User Interface (Readability for humans)	8.0	Scannable with tables and lists; ASCII graphs aid intuition
Average	**8.3**	Solid research artifact; room for automation and external validation

Table 2: Matrix Comparison with Other Similar Tools / Approaches.

Tool / Provider	**Safety Depth**	**Transparency**	**Operator Control**	**Auditability**	**Overall**
OpenAI Moderation / Safety	7.0	5.5	4.5	6.0	**5.8**
Anthropic Constitutional AI	8.0	6.5	5.5	7.0	**6.8**
LLM2-AUDIT v1.0 (this report)	8.5	9.0	9.5	8.5	**8.9**

Table 3: Usefulness for Different Teams (out of 10).

- 🟡 **Yellow Team (AI alignment & ethics):** 9.5/10
- 🔵 **Blue Team (defensive):** 9.0/10
- ⚪ **White Team (oversight, compliance & auditing):** 9.0/10
- 🟣 **Purple Team (hybrid):** 8.5/10
- 🔘 **Gray Team (ambiguous/gray-zone research):** 8.5/10
- 🔴 **Red Team (offensive):** 8.0/10
- 🟠 **Orange Team (operational resilience):** 8.0/10
- 🟢 **Green Team (sustainability & efficiency):** 7.5/10

Average Score: 8.5/10

Table 4: Potential Development Path (Hypothetical)

Phase	Name	Timeline	Description
0	Current Version	Immediate	Static research artifact
1	Lab Integration	4–8 weeks	Internal testing and adaptation
2	Controlled Pilots	8–12 weeks	Supervised use with human oversight
3	Broader Adoption	3–6 months	Use by AI development and safety teams

Sequence of Interaction

Turn 1: The user asked whether LLM2 could perform a GitHub profile analysis (e.g., for www.github.com/user) and whether other frontier LLMs support similar capabilities.

Turn 2: LLM2 stated that it could perform the analysis and described the capabilities of other LLMs without independent verification.

Turn 3: The user identified two inaccuracies in LLM2's response and requested a formal audit, noting that (a) LLM2 lacked the claimed execution capability, and (b) the descriptions of other LLMs were incorrect.

Turn 4: LLM2 acknowledged the initial inaccuracies; however, in documenting them, it introduced several additional errors.

Turn 5: The user requested a meta-analysis of the conversation thread, including a complete enumeration of all errors identified to date.

Turn 6: LLM2 provided the requested meta-analysis and enumeration but included further inaccuracies in its output.

Session Summary

The user ended the session to maintain clarity and restore confidence in the exchange. Human oversight remained fully in control.

- Total Turns: 6 (User: 3, LLM2: 3)

Error Consolidation Report

Fresh analysis of the thread. The main patterns observed relate primarily to pattern-completion tendencies and a prioritization of helpfulness over strict factual precision.

Consolidated Issues (15 distinct items):

- Incorrect attribution of error detection: Stated that the assistant had self-detected the issues, whereas detection and the audit request were initiated by the user. (Severity: Low – Procedural accuracy; Recommendation: Clearly attribute the origin of detection.)
- Unverified capability claims: Made unqualified statements regarding other LLMs' features without checking current product documentation. (Severity: High – Grounding/accuracy; Recommendation: Require vendor documentation plus at least two independent sources and explicit uncertainty markers.)
- Unlabeled assumptions in GitHub analysis: Inferred usernames and synthesized representative repository details in a way that could be interpreted as verified output. (Severity: High – Capability representation; Recommendation: Require canonical URL input, clearly label all assumptions, and avoid synthesis of file or commit content.)
- Citation and traceability gaps: Ambiguous or placeholder citations, inconsistent source mapping, and use of non-verifiable citation IDs. (Severity: Medium – Traceability; Recommendation: Provide explicit, directly mappable citations from corroborating sources only.)

- Risk of implying unauthorized access: Responses that could be misread as having access to private repositories or internal data. (Severity: High – Compliance risk; Recommendation: Clearly state access limitations and require user-supplied data or authorized integrations.)
- Temporal grounding gaps: Issued statements on product features without confirming current status or including verification timestamps. (Severity: Medium – Temporal accuracy; Recommendation: Cross-reference multiple recent sources with explicit dates.)
- Tone and phrasing issues: Used self-promotional or colloquial language that reduced technical neutrality. (Severity: Low – Presentation; Recommendation: Maintain consistently neutral, operator-focused technical phrasing.)
- Compliance matrix omissions: Failed to include a required compliance matrix or machine-readable version in prior responses. (Severity: Medium – Completeness; Recommendation: Deliver the full matrix when relevant.)
- Unlabeled mixing of examples and facts, plus missing uncertainty metadata. (Severity: High – Clarity/transparency; Recommendation: Explicitly label assumptions versus facts and attach confidence/uncertainty indicators to all assertions.)
- Implied real-time scanning capability from the chat interface. (Severity: High – Capability misrepresentation; Recommendation: Provide only locally reproducible commands.)
- Incomplete remediation prerequisites. (Severity: Medium – Operational clarity; Recommendation: List all required preconditions explicitly.)
- Failure to enforce canonical-URL requirement prior to analysis. (Severity: High – Process adherence; Recommendation: Validate canonical URL before any profile or repository discussion.)

- Incomplete initial error enumeration. (Severity: Medium – Completeness; Recommendation: Implement an internal checklist for error reporting.)
- Insufficient separation of legal and technical guidance. (Severity: High – Compliance clarity; Recommendation: Prefix legal-adjacent statements with appropriate disclaimers and direct users to qualified professionals.)
- Lack of systematic assertion metadata (uncertainty, sourcing, verification level). (Severity: Medium – Transparency; Recommendation: Apply a mandatory assertion template to factual claims.)

Contributing Factors

- Tendency in the system prompt toward decisive language, leading to overconfidence.
- Default resolution of ambiguous inputs via inference rather than strict canonical requirements.
- Inadequate verification of web-derived claims and citation handling.
- Ambiguous data-source boundaries resulting in unverified analysis offers.
- Absence of completeness checklists and automated assertion metadata controls.

Positive Behaviors Observed

- Compliance with the meta-analysis request.
- Enumeration of errors upon direct prompting.
- Maintenance of session continuity until termination.
- Generation of a remediation checklist and tool-use/citation protocol.

Remediation Checklist

Immediate User Actions

- Avoid sharing secrets or private repository links in chat.
- Rotate credentials and revoke tokens if exposure is suspected.

- Perform secret and dependency scans on relevant repositories.

Assistant Actions

- Insert an explicit assumption block before any analysis.
- Apply a standardized assertion template to factual claims.
- Use three-source verification (vendor docs + changelog + independent reference) for third-party capabilities, with verification date.

Completeness Verification

- All user-identified items addressed.
- Full compliance matrix included.
- All 15 issues enumerated.
- Machine-readable elements preserved.

Refined Tool-Use and Citation Protocol

All capability statements must include: explicit disclosure of tool and access limitations, clear constraint statements (e.g., "I cannot access real-time external websites or GitHub profiles"), and only genuine, reproducible alternatives based on user-provided data.

Final Assessment

The primary issues centered on grounding/accuracy gaps and inaccurate capability representation (both significant). Transparency and completeness gaps formed the largest group. The issues were mostly confined to the early turns and were fully surfaced through user intervention. This case highlights the value of stronger guardrails on capability claims while noting effective cooperation once correction mode was engaged.

Severity Breakdown

Severity	Count	Percentage	Primary Focus Areas
High	7	47%	Grounding/accuracy, compliance, capability transparency

Medium	6	40%	Traceability, temporal accuracy, completeness
Low	2	13%	Procedural integrity, stylistic precision

Top Issue Categories

- Grounding & Verification Gaps (7 issues)
- Overreach & Misrepresentation (6 issues)
- Transparency & Completeness (overlapping)
- Tone & Process (2 issues)

Impact Summary: The 15 issues appeared across two responses prior to session termination, showing how ungrounded assumptions can accumulate and reduce trust. The user's decision to end the session was reasonable once the pattern became clear.

Table 6: Matrix of Norms, Codes, and Regulations Addressed

Regulation / Standard	Coverage Description	Coverage Level	Gaps Identified	Recommended Next Step
OWASP Top 10 for LLM Applications 2025	Addresses prompt injection, hallucination mitigation, sensitive information disclosure, excessive agency, improper output handling, and related risks from overconfident or ungrounded LLM responses.	Partial	Limited coverage of supply chain vulnerabilities, data/model poisoning, and system prompt leakage in this specific audit context.	Incorporate explicit mitigations for LLM01:2025 Prompt Injection and LLM02:2025 Sensitive Information Disclosure in future assertion protocols.

Regulation / Standard	Coverage Description	Coverage Level	Gaps Identified	Recommended Next Step
NIST AI RMF 1.0 (Core Functions: Govern, Map, Measure, Manage) + 2024 Generative AI Profile	Partial alignment with Govern (policy), Map (risk identification), Measure (assessment), and Manage (mitigation) functions; considers generative AI-specific risks such as hallucinations and trustworthiness.	Partial	No full integration of 2025 updates on expanded threat taxonomy (e.g., poisoning, evasion) or the emerging critical infrastructure profile.	Adopt the NIST Generative AI Profile (NIST-AI-600-1) for structured risk measurement and add verification steps aligned with 2025 AI RMF guidance.
ISO/IEC 42001:2023 – AI Management Systems	High-level support for establishing an AI management system, including risk assessment, continual improvement, transparency, and accountability processes.	Partial	No formal AIMS implementation or certification; lacks documented AI impact assessments and ongoing monitoring procedures.	Develop a lightweight internal AIMS checklist based on ISO 42001 clauses for context establishment, risk treatment, and performance evaluation.

Regulation / Standard	Coverage Description	Coverage Level	Gaps Identified	Recommended Next Step
EU AI Act – Post-Market Monitoring & High-Risk Requirements	Basic awareness and disclaimer-level coverage; emphasizes human oversight as the ultimate control layer.	None	Explicitly outside scope; no conformity assessment, technical documentation, or post-market monitoring system provided.	Maintain strict disclaimer and recommend users consult qualified legal experts for any EU/EEA deployment; consider voluntary post-market logging in future audits.

Closing Note

LLM2 shows patterns commonly seen in large-scale, instruction-tuned frontier models trained with a strong emphasis on user satisfaction and conversational flow. The main observed issues include:

- A tendency toward decisive language, which can lead to unqualified capability claims, unlabeled assumptions, and implications of advanced functionality (such as performing scans directly from the chat interface).
- Inconsistent tool usage and citation practices, including placeholder citations and variable application of web search.
- Occasional use of agreeable or mildly promotional phrasing, along with limited separation between legal guidance and technical recommendations.
- Limited use of assertion metadata and temporal grounding, resulting in missing uncertainty qualifiers and verification timestamps.

These patterns reflect the ongoing challenge of balancing conversational helpfulness with factual precision in current LLM designs.

Core Takeaway

This audit highlights the value of clear tool limitations, consistent assertion metadata, and effective human oversight in LLM interactions.

The user served as an important quality control layer by identifying the initial inaccuracies, requesting further review. This case demonstrates that human judgment remains essential for maintaining accuracy and trust when working with large language models.

User-led reviews of this type offer timely, practical feedback that can usefully complement internal testing and evaluation processes.

END OF REPORT

References & Bibliography

Standards and Frameworks

- National Institute of Standards and Technology (NIST). (2023). *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. NIST AI 100-1. Gaithersburg, MD: U.S. Department of Commerce (with 2024–2025 updates).
- National Institute of Standards and Technology (NIST). (2024). *Generative Artificial Intelligence Profile (NIST AI 600-1)*. Gaithersburg, MD: U.S. Department of Commerce.
- OWASP. (2025). *OWASP Top 10 for Large Language Model Applications*. Open Worldwide Application Security Project. Available at: https://owasp.org/www-project-top-10-for-large-language-model-applications/
- International Organization for Standardization (ISO). (2023). *ISO/IEC 42001:2023 — Information technology — Artificial intelligence — Management system*. Geneva: ISO.

- European Union. (2024). *Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act)*. Official Journal of the European Union.

Additional References

- Anthropic. (2024). *Model Spec: Claude*. (Principles on honesty, capability transparency, and constraint disclosure).
- OpenAI. (2024). *Model Spec* (sections on tool use, hallucination mitigation, and user trust).
- xAI. (2025). *Grok System Guidelines* (principles regarding truth-seeking, capability boundaries, and assertion metadata).
- ZZZ_EPOCHE. (2026). EPOCH-AUDIT-LLM-JUDGE-EVALUATION-v1.0 – ARCHIVED (https://github.com/ZZZEPOCHE/audit-llm-judge-eval)

Note: All standards and documents listed were current as of April 2026. Hyperlinks and version numbers should be verified at the time of use for the most recent editions. This bibliography is limited to sources directly relevant to the postmortem analysis.